
Multicore proj3-prob2



subject	multicore
professor	손봉수 교수님
date	2023.05.20
department	소프트웨어대학 소프트웨어학부
name	20200641 임수현

Environment

CPU : 11th Gen Intel Core(TM) i7-1195G7 //Quard Core
Core : 4
RAM : 16GB
System type : 64 bits
OS : Window 11
Run in WSL

Source Code

```
#include <stdlib.h>
#include <stdbool.h>
#include <omp.h>
#include <stdio.h>

long num_steps = 100000000;
double step;

void main (int argc, char *argv[]){
    long i; double x, pi, sum = 0.0;
    double start_time, end_time;

    int schedul_type = atoi(argv[1]);
    int chunk_size=atoi(argv[2]);
    int thread_num=atoi(argv[3]);

    step = 1.0/(double) num_steps;

    switch (schedul_type){
    case 1:
        start_time = omp_get_wtime();
        omp_set_num_threads(thread_num);
        {
            #pragma omp parallel for reduction(+:sum) private(x)
            schedule(static,chunk_size)
                for (i=0;i< num_steps; i++){
                    x = (i+0.5)*step;
                    sum = sum + 4.0/(1.0+x*x);
                }
        }
        break;
    case 2:
        start_time = omp_get_wtime();
        omp_set_num_threads(thread_num);
        {
            #pragma omp parallel for reduction(+:sum) private(x)
            schedule(dynamic,chunk_size)
                for (i=0;i< num_steps; i++){
                    x = (i+0.5)*step;
                    sum = sum + 4.0/(1.0+x*x);
                }
        }
        break;
    }
```

```

case 3:
    start_time = omp_get_wtime();
    omp_set_num_threads(thread_num);
    {
        #pragma omp parallel for reduction(+:sum) private(x)
        schedule(guided,chunk_size)
            for (i=0;i< num_steps; i++){
                x = (i+0.5)*step;
                sum = sum + 4.0/(1.0+x*x);
            }
        break;
    }
    pi = step * sum;
    end_time = omp_get_wtime();

    double timeDiff = end_time - start_time;
    printf("Execution Time : %lfms\n", timeDiff*1000);

    printf("pi=%.24lf\n",pi);
}

```

code explain

After receiving the schedule type, each schedule type was executed using a switch statement. The number of threads inputted in the *omp_set_num_thread* was put as a parameter. In addition, the pi value was obtained by calculating through a for statement using *#pragma omp parallel for reduction (+:sum) private(x) schedule (schedule type, chunk_size)*. The sum was set to *public* because all threads must be able to refer to it, and the x value to be used for each thread to calculate was set to *private*.

Result



execution time (msec)	chunk size	1	2	4	6
static	1	88.07916	44.98113	23.08272	15.64471
dynamic		141.6732	142.619	227.9872	214.9876
guided		91.88717	44.55762	24.54581	16.08232
static	5	90.13424	45.18829	22.93662	15.79019
dynamic		102.129	96.67987	62.77164	54.1078
guided		88.76649	44.59407	22.79211	15.51427
static	10	88.7931	44.67053	23.25887	15.92089
dynamic		94.12386	75.12677	39.7092	30.26232
guided		88.52559	44.45776	22.78815	15.54352
static	100	91.75869	44.80331	22.84745	15.60365
dynamic		88.43275	46.09615	24.54836	16.08396
guided		88.11994	44.70774	22.89749	15.77961

execution time (msec)	chunk size	8	10	12	14	16
static	1	12.8087	18.2988	16.4638	14.13903	18.00131
dynamic		195.0464	204.5358	195.0464	194.0668	195.2742
guided		12.69214	12.08618	12.15801	12.15558	12.36472
static	5	16.87883	19.01885	15.67211	13.86019	17.6868
dynamic		45.75027	45.50876	45.55088	45.89917	46.59605
guided		11.88803	12.02698	12.23279	11.96602	13.40366
static	10	11.79512	18.81129	15.92299	13.64646	12.88441
dynamic		23.93501	24.50997	24.60219	24.6474	24.52228
guided		11.96801	12.49555	12.04673	13.14405	12.40213
static	100	11.85371	18.5047	15.92591	13.85421	12.40443
dynamic		12.05684	12.54629	14.44072	12.55766	12.56636
guided		12.1033	12.30475	12.25909	14.13145	12.56208

performance (1/exec time)	chunk size	1	2	4	6
static	1	0.011353	0.022232	0.043322	0.063919
dynamic		0.007059	0.007012	0.004386	0.004651
guided		0.010883	0.022443	0.04074	0.06218
static	5	0.011095	0.02213	0.043598	0.06333
dynamic		0.009792	0.010343	0.015931	0.018482
guided		0.011266	0.022425	0.043875	0.064457
static	10	0.011262	0.022386	0.042994	0.062811
dynamic		0.010624	0.013311	0.025183	0.033044
guided		0.011296	0.022493	0.043882	0.064335
static	100	0.010898	0.02232	0.043769	0.064088
dynamic		0.011308	0.021694	0.040736	0.062174
guided		0.011348	0.022367	0.043673	0.063373



performance (1/exec time)	chunk size	8	10	12	14	16
static	1	0.078072	0.054648	0.060739	0.070726	0.055552
dynamic		0.005127	0.004889	0.005127	0.005153	0.005121
guided		0.078789	0.082739	0.08225	0.082267	0.080875
static	5	0.059246	0.052579	0.063808	0.072149	0.056539
dynamic		0.021858	0.021974	0.021953	0.021787	0.021461
guided		0.084118	0.083146	0.081748	0.08357	0.074606
static	10	0.084781	0.05316	0.062802	0.073279	0.077613
dynamic		0.04178	0.0408	0.040647	0.040572	0.040779
guided		0.083556	0.080028	0.08301	0.07608	0.080631
static	100	0.084362	0.05404	0.062791	0.07218	0.080616
dynamic		0.08294	0.079705	0.069249	0.079633	0.079578
guided		0.082622	0.081269	0.081572	0.070764	0.079605

Static scheduling





The performance is similar regardless of the chunk size because the calculation inside the for statement calculates a similar amount for each thread chunk. It can be seen that the workload is evenly distributed, so that load balancing between threads is maintained to some extent. And as the number of threads increases, the performance improves and then becomes constant. I think this is due to the occurrence of communication overhead. Therefore, increasing the amount of threads does not necessarily improve performance.

Dynamic scheduling



It can be seen that the performance is different depending on the chunk size. Smaller chunk size can cause overhead in allocating and adjusting tasks. So it can be seen that the smaller the chunk size, the lower the performance. And as the number of threads increases, the performance improves, and if there are more than 8, the performance does not increase. Like static, it can be seen that communication overhead occurs and performance is no longer increased.

Guided scheduling



The guided method first starts with a large chunk size and then decreases the chunk size as it goes back, helping the thread that finished the calculation quickly to make the next calculation. Therefore, regardless of the chunk size, the exec time and performance result are the same. Even if the chunk size is small, all threads are busy at the beginning of the run, so the chunk size doesn't matter. However, the chunk size becomes smaller as it goes back, and it can be seen that the performance improves as the number of threads processed by the calculated thread increases. However, performance does not continue to increase. When the number of threads increases and the number of threads exceeds 8, the exec time is no longer reduced. This means that there is no difference in performance even if the number of threads increases due to communication overhead between threads.